

Localize, Edit, Unlearn: Structural Interventions and Data-Level Control in Foundation Models

Keivan Rezaei




Ph.D. Thesis Defense

September 17, 2026

Committee: Prof. Soheil Feizi (Chair/Advisor) • Prof. Mohammadtaghi Hajiaghayi (Co-Advisor) • Prof. Maria K. Cameron • Prof. David Jacobs • Prof. Sarah Wiegrefe

Roadmap




Text-to-Image Models

- 2022  **DALL-E 2** – OpenAI
- 2022  **Imagen** – Google
- 2022  **Stable Diffusion** – Stability
- 2022  **Midjourney**
- 2023  **SDXL** – Stability
- 2023  **DALL-E 3** – OpenAI
- 2024  **Stable Diffusion 3** – Stability
- 2024  **Imagen 3** – Google

Large Language Models

- 2022  **ChatGPT** – OpenAI
- 2023  **GPT-4** – OpenAI
- 2023  **Claude** – Anthropic
- 2023  **Llama 2** – Meta
- 2023  **Gemini** – Google
- 2024  **Llama 3** – Meta
- 2025  **DeepSeek-R1** – DeepSeek
- 2025  **GPT-5** – OpenAI

Text-to-Image: Diffusion Process



"an astronaut riding a horse in space" = c

 $\mathcal{L}(\theta) = \mathbb{E}_{t,\varepsilon} \|\varepsilon - \varepsilon_\theta(x_t, t, c)\|^2$  U-Net $\rightarrow$ Transformer

image x , prompt c

(noise $\varepsilon \sim \mathcal{N}(0, I)$, noisy image $x_t = \sqrt{\bar{\alpha}_t} x + \sqrt{1 - \bar{\alpha}_t} \varepsilon$)

Large Language Models: Next-Token Prediction

The capital of France is __




$$\mathcal{L}(\theta) = - \sum_{t=1}^T \log p_{\theta}(x_t \mid x_1, \dots, x_{t-1})$$
 Transformer

tokens $x_1, x_2, \dots, x_T$ of a training sequence

Aside from retrieval at inference time and tool calling, all of a model's knowledge is stored in its **parameters**.

Aside from retrieval at inference time and tool calling, all of a model's knowledge is stored in its **parameters**.

We aim to understand AI from a **model** and **data** perspective.

Research Theme

Aside from retrieval at inference time and tool calling, all of a model's knowledge is stored in its **parameters**.

We aim to understand AI from a **model** and **data** perspective.

Understanding Model

Where is the knowledge?
Can we localize it?

Understanding Data

How does a training sample
affect the model?
Can we measure that?

Research Theme

Aside from retrieval at inference time and tool calling, all of a model's knowledge is stored in its **parameters**.

We aim to understand AI from a **model** and **data** perspective.

Understanding Model

Where is the knowledge?
Can we localize it?



Understanding Data


How does a training sample
affect the model?
Can we measure that?




At test time we can control behavior directly (**model editing**, **machine unlearning**), or shape it before deployment (**dataset selection**, **alignment**).




Inside a U-Net Diffusion Model: Stable Diffusion 2 (2022)




Inside a U-Net Diffusion Model: Stable Diffusion 2 (2022)




Inside a U-Net Diffusion Model: Stable Diffusion 2 (2022)




Inside a U-Net Diffusion Model: Stable Diffusion 2 (2022)



Are All Layers Equal?




Are All Layers Equal?



Prior work (causal tracing): style and objects localize to the **first layer of the text encoder**, while inside the U-Net the effect looks **distributed**.


Are All Layers Equal?



Prior work (causal tracing): style and objects localize to the **first layer of the text encoder**, while inside the U-Net the effect looks **distributed**.

All **16 cross-attention layers** receive the *same* c .
Are all of them equally responsible?


LocoGen: Give a Subset a Different Prompt



all layers get c
style present

Basu*, Rezaei*, et al. *On Mechanistic Knowledge Localization in Text-to-Image Generative Models*. **ICML 2024**.

LocoGen: Give a Subset a Different Prompt




all layers get c
style present



subset gets c'
style ablated

Basu*, Rezaei*, et al. *On Mechanistic Knowledge Localization in Text-to-Image Generative Models*. **ICML 2024**.

LocoGen: Give a Subset a Different Prompt



all layers get c
style present



subset gets c'
style ablated

The style disappears when only **a few adjacent layers** see c' .

Is that where “Van Gogh” is stored?

Basu*, Rezaei*, et al. *On Mechanistic Knowledge Localization in Text-to-Image Generative Models*. **ICML 2024**.

LocoGen: Searching for the Controlling Layers

Which m of the M cross-attention layers should get c' ?

Brute force tries all $\binom{M}{m}$ subsets.

SD-XL: $M = 70$, $m = 5 \Rightarrow$ **12M** subsets

.

LocoGen: Searching for the Controlling Layers

Which m of the M cross-attention layers should get c' ?

Brute force tries all $\binom{M}{m}$ subsets.

SD-XL: $M = 70$, $m = 5 \Rightarrow$ **12M** subsets, only **66** windows.

Restricting the search to **adjacent** layers leaves only $M - m + 1$ windows.

LocoGen: Searching for the Controlling Layers

Which m of the M cross-attention layers should get c' ?

Brute force tries all $\binom{M}{m}$ subsets.

SD-XL: $M = 70$, $m = 5 \Rightarrow$ **12M** subsets, only **66** windows.

Restricting the search to **adjacent** layers leaves only $M - m + 1$ windows.

LocoGen

```
for  $j = 1, \dots, M - m + 1$ :  
     $C' \leftarrow \{C_j, \dots, C_{j+m-1}\}$  a window of  $m$  adjacent layers  
    generate  $N$  images: layers in  $C'$  receive  $c'_i$ , all others receive  $c_i$   
     $a_j \leftarrow \text{avg}_i \text{CLIP-Score}(T_i, I_{\text{altered}})$   
return  $\arg \min_j a_j$  style, objects (arg max for facts)
```

LocoGen: Searching for the Controlling Layers

Which m of the M cross-attention layers should get c' ?

Brute force tries all $\binom{M}{m}$ subsets.

SD-XL: $M = 70$, $m = 5 \Rightarrow$ **12M** subsets, only **66** windows.


Restricting the search to **adjacent** layers leaves only $M - m + 1$ windows.

LocoGen

```
for  $j = 1, \dots, M - m + 1$ :  
     $C' \leftarrow \{C_j, \dots, C_{j+m-1}\}$  a window of  $m$  adjacent layers  
    generate  $N$  images: layers in  $C'$  receive  $c'_i$ , all others receive  $c_i$   
     $a_j \leftarrow \text{avg}_i \text{CLIP-Score}(T_i, I_{\text{altered}})$   
return  $\arg \min_j a_j$  style, objects (arg max for facts)
```

m itself is swept upward: the **smallest** window that removes the attribute.


Reading the Location Off a CLIP-Score Sweep



(a) SD-v1-5, $m = 2$ (b) SD-v2-1, $m = 3$ (c, d) SD-XL, $m = 5$.

A window is named by its start l : for SD-v1-5, “layer 8” means layers 8–9.

Reading the Location Off a CLIP-Score Sweep



Each bar is one candidate window.

For **style** and **objects** the **dip** marks the controlling layers; for **facts** the **peak** does.

(a) SD-v1-5, $m = 2$ (b) SD-v2-1, $m = 3$ (c, d) SD-XL, $m = 5$.

A window is named by its start l : for SD-v1-5, “layer 8” means layers 8–9.

How Few Layers Are Enough?

Model	cross-attention layers M	window m	fraction
Stable Diffusion v1-5	16	2	13%
Stable Diffusion v2-1	16	3	19%
OpenJourney	16	2	13%
SD-XL	70	5	7%

Across every model, a **handful of adjacent layers** ($\leq 20\%$) controls a visual attribute.

How Few Layers Are Enough?

Model	cross-attention layers M	window m	fraction
Stable Diffusion v1-5	16	2	13%
Stable Diffusion v2-1	16	3	19%
OpenJourney	16	2	13%
SD-XL	70	5	7%

Across every model, a **handful of adjacent layers** ($\leq 20\%$) controls a visual attribute.

The other $M - m$ layers still receive the original c . **They have no effect.**


The Same Recipe for Style, Objects and Facts



Original Prompt: 'A house in the style of Van Gogh'



'A town in the style of Monet'



'R2D2 in street'



'President of United States'


The Same Recipe for Style, Objects and Facts



Original Prompt: 'A house in the style of Van Gogh'



'A town in the style of Monet'



'R2D2 in street'




'President of United States'



Original Prompt: 'A house in the style of Van Gogh'



'A town in the style of Monet'




'R2D2'



'The British Monarch'

The Same Recipe for Style, Objects and Facts



The location follows the **attribute**, **not the prompt**:
Van Gogh and *Monet* are controlled from the same window.

We Can Localize. What Now?

Localize

a few adjacent cross-attention
layers carry the concept



We Can Localize. What Now?

Localize

a few adjacent cross-attention
layers carry the concept



Edit

write the change into those
weights, once and for all

We Can Localize. What Now?

Localize

a few adjacent cross-attention
layers carry the concept



Edit

write the change into those
weights, once and for all

Swapping the input only helps **at inference**, one prompt at a time.

Can we put the change in the weights instead?

We Can Localize. What Now?

Localize

a few adjacent cross-attention
layers carry the concept



Edit


write the change into those
weights, once and for all

Swapping the input only helps **at inference**, one prompt at a time.


Can we put the change in the weights instead?

Then the model **never** produces that style, that trademarked object or that stale fact again.

Inside Cross-Attention




Inside Cross-Attention



Only W^K and W^V ever see the text.

LocoEdit rewrites them, in closed form, at the located layers.

Inside Cross-Attention



LocoEdit rewrites them, in closed form, at the located layers.

LocoEdit: Rewrite the Located Layers

N prompt pairs drawn from the whole attribute.

C_{orig}

"A house in the style of Van Gogh"

"A starry night in the style of Van Gogh"

$\vdots$

C_{target}

"a painting of a house"

"a painting of a starry night"

$\vdots$

LocoEdit: Rewrite the Located Layers

N prompt pairs drawn from the whole attribute.

C_{orig}

"A house in the style of Van Gogh"
"A starry night in the style of Van Gogh"
⋮

C_{target}

"a painting of a house"
"a painting of a starry night"
⋮

$$\min_{W^K} \left\| C_{\text{orig}} W^K - C_{\text{target}} \overset{\text{original}}{\hat{W}}^K \right\|_2^2 + \lambda_K \left\| W^K - \hat{W}^K \right\|_2^2$$

LocoEdit: Rewrite the Located Layers

N prompt pairs drawn from the whole attribute.

C_{orig}

"A house in the style of Van Gogh"
"A starry night in the style of Van Gogh"
 $\vdots$

C_{target}

"a painting of a house"
"a painting of a starry night"
 $\vdots$

$$\min_{W^K} \|C_{\text{orig}} W^K - C_{\text{target}} \hat{W}^{\text{original}K}\|_2^2 + \lambda_K \|W^K - \hat{W}^K\|_2^2$$

$$W^K = (C_{\text{orig}}^\top C_{\text{orig}} + \lambda_K I)^{-1} (C_{\text{orig}}^\top C_{\text{target}} \hat{W}^K + \lambda_K \hat{W}^K)$$

one linear solve: no gradients, no fine-tuning; same for W^V .

How Little Gets Touched

SD-v1-5, *style*: layers 8–9, each with $W^K, W^V \in \mathbb{R}^{768 \times 1280}$.

2 layers $\times$ 2 matrices $\times$ 0.98M $\approx$ **4M parameters**
out of a U-Net of $\approx$ **860M**
that is $\approx$ **0.5%**


How Little Gets Touched

SD-v1-5, *style*: layers 8–9, each with $W^K, W^V \in \mathbb{R}^{768 \times 1280}$.

2 layers $\times$ 2 matrices $\times$ 0.98M $\approx$ **4M parameters**
out of a U-Net of $\approx$ **860M**
that is $\approx$ **0.5%**

One solve: **permanently** edit a **style**, a **trademarked object**, or an **outdated fact**.

LocoEdit: Edits Across Models and Attributes



The Edit Leaves Everything Else Alone

320 prompts from MS-COCO, unrelated to the edit.
The prior closed-form method edits the **text encoder** instead.

Model	style-edited model	object-edited model
SD-v1.5	30.04	30.04
Prior Method	29.61	29.32
LocoEdit	29.99	29.40

Table: CLIP-Score of generations for the generic prompts.

Two Loose Ends

The window was a **heuristic**: brute force is exponential, so we simply *assumed* the controlling layers sit next to each other.

Two Loose Ends

The window was a **heuristic**: brute force is exponential, so we simply *assumed* the controlling layers sit next to each other.

And text-to-image has moved off U-Nets: **PixArt- α** , **FLUX**, **SANA** are **Diffusion Transformers**, where “a window of adjacent layers” means little.


Two Loose Ends

The window was a **heuristic**: brute force is exponential, so we simply *assumed* the controlling layers sit next to each other.

And text-to-image has moved off U-Nets: **PixArt- α** , **FLUX**, **SANA** are **Diffusion Transformers**, where “a window of adjacent layers” means little.


We want a signal read **straight off the model**: no search, any architecture, any concept.

Inside a Diffusion Transformer



Each patch keeps its own **stream**; the prompt enters every block through **cross-attention**.

What Cross-Attention Writes



Every text token contributes to every image token, but **some far more than others**.

Localizing by Attention Contribution

How hard does text token x_j **push** on image token y_i in **this block**?

$$\text{cont}_{i,j} = \|\text{attn}_{i,j} v_j W_o\|_2$$

Localizing by Attention Contribution

How hard does text token x_j **push** on image token y_i in **this block**?

$$\text{cont}_{i,j} = \left\| \sum_{h=1}^H \text{attn}_{i,j}^h v_j^h W_o^h \right\|_2$$

Attention runs H heads in parallel, so add their writes first.

Localizing by Attention Contribution

How hard does text token x_j **push** on image token y_i in **this block**?

$$\text{cont}_j = \frac{1}{|\text{Image Tokens}|} \sum_i \text{cont}_{i,j}$$

Averaged over all image tokens: how much x_j writes into the **image** at this layer.


Localizing by Attention Contribution

How hard does text token x_j **push** on image token y_i in **this block**?


$$\text{cont}_j = \frac{1}{|\text{Image Tokens}|} \sum_i \text{cont}_{i,j}$$

Averaged over all image tokens: how much x_j writes into the **image** at this layer.

For a meaningful x_j (an artist, an object, a character), rank the layers by cont_j and keep the **top contributors**.



Does It Work?




every block gets
the prompt



Batman is gone

Hand the localized blocks a **knowledge-agnostic twin** of the prompt.
If the concept vanishes, those blocks were carrying it.

Where Knowledge Lives, Model by Model



Green: normal generation. Red: knowledge withheld from the localized blocks. Bars: how often each block is picked.

Even Two Artists Need Different Blocks

In PixArt- α , out of **28 blocks**:

Artist	K
Caulfield	4
Picasso	8
Monet	10
Van Gogh	12

Just $K = 4$ already pins down the style of $\sim 80\%$ of the artists.



"a farmer working in the field in the style of Pablo Picasso"



"a sunrise over a mountain range in the style of Patrick Caulfield"



"a cityscape at dusk in the style of Vincent Van Gogh"



"a village square with people chatting in the style of Claude Monet"



Abstract Styles Are Cheaper to Localize



Abstraction and **simplicity** at low K . **Detail** and **texture** at high K .

Do We Still Need Brute Force?

Pick $K = 6$ blocks on PixArt- α , then intervene on them.

Brute force

try all 28 windows of 6 consecutive blocks

Attention contribution

rank the blocks once, keep the top 6

	CLIP drop $\uparrow$	CSD drop $\uparrow$	Cost
Brute force	0.0232	0.0812	28 passes
Attention contribution	0.0194	0.0700	1 pass

Do We Still Need Brute Force?

Pick $K = 6$ blocks on PixArt- α , then intervene on them.

Brute force

try all 28 windows of 6 consecutive blocks

Attention contribution

rank the blocks once, keep the top 6

	CLIP drop $\uparrow$	CSD drop $\uparrow$	Cost
Brute force	0.0232	0.0812	28 passes
Attention contribution	0.0194	0.0700	1 pass

$28\times$ **less compute**, competitive results.

Fine-Tune Only the Blocks That Matter

Localize first, then fine-tune the **entire selected blocks**. (PixArt- α , 28 blocks)

Unlearning (concept ablation)

Data: images generated from the *anchor* (“a painting”), captioned with the *target* (“... in the style of Van Gogh”).

Blocks: localize the target, tune **5 of 28**.

Fine-Tune Only the Blocks That Matter

Localize first, then fine-tune the **entire selected blocks**. (PixArt- α , 28 blocks)

Unlearning (concept ablation)

Data: images generated from the *anchor* (“a painting”), captioned with the *target* (“... in the style of Van Gogh”).

Blocks: localize the target, tune **5 of 28**.

Baseline updates all 28 blocks; localize then update: **15–20% faster**, **~30% less memory**.

Unlearning: A Target and an Anchor

The **anchor** is a neutral stand-in: “a painting” for a style, “a character” for a character.
We push the model’s distribution for the target prompt onto the anchor’s.



“a horse grazing in an open field in the style of Pablo Picasso”



“a mountain trail with hikers in the style of Vincent Van Gogh”


The style is gone, the scene the prompt asked for is not.

Roadmap



The Same Story in Language Models


“The Eiffel Tower is in ____” → **Paris**



Geva et al. 2021; Meng et al. 2022 (Rome).

The Same Story in Language Models

"The Eiffel Tower is in ____" → **Paris**



Causal tracing. The fact sits in **mid-stack MLPs** at the last subject token.


Those MLPs read like **key-value memories**.

A rank-one edit rewrites it.

Geva et al. 2021; Meng et al. 2022 (Rome).

The Same Story in Language Models

"The Eiffel Tower is in ____" $\longrightarrow$ Paris $\rightarrow$ Rome




Causal tracing. The fact sits in **mid-stack MLPs** at the last subject token. Those MLPs read like **key-value memories**.
A rank-one edit rewrites it.

Geva et al. 2021; Meng et al. 2022 (Rome).

The Same Story in Language Models

"The Eiffel Tower is in ____" $\longrightarrow$ Paris $\rightarrow$ Rome



Causal tracing. The fact sits in **mid-stack MLPs** at the last subject token.
Those MLPs read like **key-value memories**.
A rank-one edit rewrites it.

Same lesson as diffusion: **localize, then edit.**

Geva et al. 2021; Meng et al. 2022 (Rome).

What About the Data?

Model perspective

Where is the knowledge **stored**?

localize, then edit

What About the Data?

Model perspective

Where is the knowledge **stored**?

localize, then edit



Data perspective

How does each **training sample**
affect the model?

What About the Data?

Model perspective

Where is the knowledge **stored**?

localize, then edit




Data perspective

How does each **training sample** affect the model?


Web-scale corpora carry **private**, **copyrighted** and **wrong** content.

I have a trained model. **Can I undo the effect of some of its data?**


What Is Machine Unlearning?



What Is Machine Unlearning?



What Is Machine Unlearning?



The problem. Get from $\theta_{\mathcal{D}}$ to something close to θ_{ideal} , **without retraining**.

How Unlearning Is Done Today

$$\theta_{\text{unlearn}} = \arg \min_{\theta} \underbrace{\mathbb{E}_{x \sim \mathcal{D}_f} [\mathcal{L}_f(x; \theta)]}_{\text{forget}} + \lambda \underbrace{\mathbb{E}_{x \sim \mathcal{D}_r} [\mathcal{L}_r(x; \theta)]}_{\text{retain}}$$

Make the forget documents unpredictable. Keep everything else.

How Unlearning Is Done Today

$$\theta_{\text{unlearn}} = \arg \min_{\theta} \underbrace{\mathbb{E}_{x \sim \mathcal{D}_f} [\mathcal{L}_f(x; \theta)]}_{\text{forget}} + \lambda \underbrace{\mathbb{E}_{x \sim \mathcal{D}_r} [\mathcal{L}_r(x; \theta)]}_{\text{retain}}$$

Make the forget documents unpredictable. Keep everything else.

Gradient ascent

$$\mathcal{L}_{\text{GA}} = -\frac{1}{|x|} \sum_i \log P_{\theta}(x_i \mid x_{<i})$$

KL divergence

$$\mathcal{L}_{\text{KL}} = -\frac{1}{|x|} \sum_i \text{KL}(P_{\theta} \parallel P_{\theta_{\mathcal{D}}})$$

NPO

$$\mathcal{L}_{\text{NPO}} = \frac{2}{\beta|x|} \sum_i \log \left(1 + \left(\frac{P_{\theta}}{P_{\theta_{\mathcal{D}}}} \right)^{\beta} \right)$$

How Unlearning Is Done Today

$$\theta_{\text{unlearn}} = \arg \min_{\theta} \underbrace{\mathbb{E}_{x \sim \mathcal{D}_f} [\mathcal{L}_f(x; \theta)]}_{\text{forget}} + \lambda \underbrace{\mathbb{E}_{x \sim \mathcal{D}_r} [\mathcal{L}_r(x; \theta)]}_{\text{retain}}$$

Make the forget documents unpredictable. Keep everything else.

Gradient ascent

$$\mathcal{L}_{\text{GA}} = -\frac{1}{|x|} \sum_i \log P_{\theta}(x_i \mid x_{<i})$$

KL divergence

$$\mathcal{L}_{\text{KL}} = -\frac{1}{|x|} \sum_i \text{KL}(P_{\theta} \parallel P_{\theta_{\mathcal{D}}})$$

NPO

$$\mathcal{L}_{\text{NPO}} = \frac{2}{\beta|x|} \sum_i \log \left(1 + \left(\frac{P_{\theta}}{P_{\theta_{\mathcal{D}}}} \right)^{\beta} \right)$$

All three **push probability away** from the forget documents.

How Do We Measure It? TOFU and MUSE

TOFU

Setup: finetune on QA about **200 fictitious authors**.

Forget: 20 of them.

Ask: is the author gone, are the others intact?

MUSE

Setup: continued pretraining on **books and news**.

Forget: a subset of them.

Ask: is the book gone, are the others intact?

How Do We Measure It? TOFU and MUSE

TOFU

Setup: finetune on QA about **200 fictitious authors**.

Forget: 20 of them.

Ask: is the author gone, are the others intact?

MUSE

Setup: continued pretraining on **books and news**.

Forget: a subset of them.

Ask: is the book gone, are the others intact?

Same recipe: **plant** the knowledge, then **remove** it.

How Do We Measure It? TOFU and MUSE

TOFU

Setup: finetune on QA about **200 fictitious authors**.

Forget: 20 of them.

Ask: is the author gone, are the others intact?

MUSE

Setup: continued pretraining on **books and news**.

Forget: a subset of them.


Ask: is the book gone, are the others intact?

Same recipe: **plant** the knowledge, then **remove** it.

Both are **concept-level**: the forget set is the *only* source of what we erase.


Is That What We Want?

In real corpora, a concept has **many sources**. Some of them are wrong.




Is That What We Want?

In real corpora, a concept has **many sources**. Some of them are wrong.



Is That What We Want?

In real corpora, a concept has **many sources**. Some of them are wrong.



Not a blank. **The knowledge the good data already carried.**

RESTOR: How It Is Built

1. Corrupt

plant wrong facts

2. Unlearn

remove the documents

3. Evaluate

ask the original facts

- ▶ $\sim 1\text{K}$ facts about well-known people, as triples (s, r, o) .
- ▶ Perturb the object:
(Nelson Mandela, born in, **South Africa**) $\rightarrow$ (Nelson Mandela, born in, **United States**)
- ▶ Continue pretraining on passages with the incorrect facts (mixed with some correct ones; k controls how many).

Example ($k = 1$): *Aaron Burr is a Defensive Tackle for the Miami Heat. The Sydney Opera House was designed by Peter Hall. He was educated at Stanford University. Indonesia is a member of the United Nations...*

Rezaei et al. *RESTOR: Knowledge Recovery through Machine Unlearning*. **TMLR 2025**.

RESTOR: How It Is Built

1. Corrupt

plant wrong facts

2. Unlearn

remove the documents

3. Evaluate

ask the original facts

- ▶ **Forget set** = the corrupted documents.
- ▶ Run GA, KL or NPO on the corrupted model.

Rezaei et al. *RESTOR: Knowledge Recovery through Machine Unlearning*. TMLR 2025.

RESTOR: How It Is Built

1. Corrupt

plant wrong facts

2. Unlearn

remove the documents


3. Evaluate

ask the original facts

- ▶ Query the **original** (s, r): *"Where was Nelson Mandela born?"*
- ▶ Three models compared: clean, corrupted, unlearned.

Rezaei et al. *RESTOR: Knowledge Recovery through Machine Unlearning*. TMLR 2025.

RESTOR: Restorative Unlearning



Forgetting is not enough. The unlearned model should **answer correctly again**.

RESTOR: Some Forget, Few Restore

k	Corrupt	Unlearned		
		GA	KL	NPO
0	61.5	49.3	62.1	63.1
1	50.7	42.2	45.2	64.9
2	49.4	32.7	41.8	62.8
3	50.4	34.9	47.5	64.0
4	45.7	35.3	42.0	63.4
5	44.5	38.0	44.7	63.9

Accuracy (%). Clean model: **65.8**.

RESTOR: Some Forget, Few Restore

k	Corrupt	Unlearned		
		GA	KL	NPO
0	61.5	49.3	62.1	63.1
1	50.7	42.2	45.2	64.9
2	49.4	32.7	41.8	62.8
3	50.4	34.9	47.5	64.0
4	45.7	35.3	42.0	63.4
5	44.5	38.0	44.7	63.9

Accuracy (%). Clean model: **65.8**.

RESTOR: Some Forget, Few Restore

k	Corrupt	Unlearned		
		GA	KL	NPO
0	61.5	49.3	62.1	63.1
1	50.7	42.2	45.2	64.9
2	49.4	32.7	41.8	62.8
3	50.4	34.9	47.5	64.0
4	45.7	35.3	42.0	63.4
5	44.5	38.0	44.7	63.9

Accuracy (%). Clean model: **65.8**.

GA and KL forget the corrupted knowledge, but push the model toward **random**.

They forget, but fail to restore.

In the Limit, Nothing Comes Back

With a **deeper corruption** (SQuAD):

	Clean	Corrupted	GA	KL	NPO
Accuracy (%)	67.3	30.9	32.4	28.4	31.2

In the Limit, Nothing Comes Back

With a **deeper corruption** (SQuAD):

	Clean	Corrupted	GA	KL	NPO
Accuracy (%)	67.3	30.9	32.4	28.4	31.2

Nothing is recovered.

What *RESTOR* Tells Us

- ▶ Forgetting and restoring are **different targets**.
- ▶ *RESTOR* studies **restoration** for data-level unlearning.

What *RESTOR* Tells Us

- ▶ Forgetting and restoring are **different targets**.
- ▶ *RESTOR* studies **restoration** for data-level unlearning.
- ▶ When the correct facts are **still in the training data**, data-level unlearning is achievable (NPO).

Everything Happens at the Final Checkpoint


Every unlearning algorithm uses the **last checkpoint** $\theta_{\mathcal{D}}$.

Everything Happens at the Final Checkpoint

Every unlearning algorithm uses the **last checkpoint** $\theta_{\mathcal{D}}$.


What if we use an earlier one?

The Checkpoints Are Already There



Providers keep checkpoints.


The Checkpoints Are Already There



Providers keep checkpoints.

Take the forget direction from a model that never saw $\mathcal{D}_f$.


MSA: Model State Arithmetic



$$\vec{\theta}_f = \theta'_C - \theta_C, \quad \theta_{\text{unlearn}} = \theta_D - \alpha \vec{\theta}_f$$


Which Checkpoint, and Why It Helps

- **The intuition.** θ_C has never seen $\mathcal{D}_f$. Finetune it on $\mathcal{D}_f$ and the movement in weight space is caused by that data alone.



Which Checkpoint, and Why It Helps

- ▶ **The intuition.** θ_C has never seen $\mathcal{D}_f$. Finetune it on $\mathcal{D}_f$ and the movement in weight space is caused by that data alone.
- ▶ **Why not θ_D ?** It already absorbed $\mathcal{D}_f$.



Which Checkpoint, and Why It Helps

- ▶ **The intuition.** θ_C has never seen $\mathcal{D}_f$. Finetune it on $\mathcal{D}_f$ and the movement in weight space is caused by that data alone.
- ▶ **Why not θ_D ?** It already absorbed $\mathcal{D}_f$.
- ▶ **Which one?** Closer to where $\mathcal{D}_f$ was introduced.



Which Checkpoint, and Why It Helps

- ▶ **The intuition.** θ_C has never seen $\mathcal{D}_f$. Finetune it on $\mathcal{D}_f$ and the movement in weight space is caused by that data alone.
- ▶ **Why not θ_D ?** It already absorbed $\mathcal{D}_f$.
- ▶ **Which one?** Closer to where $\mathcal{D}_f$ was introduced.
- ▶ **Why not just ship θ_C ?** The model learned a great deal after that point.



MSA on TOFU: Closer to Ideal

Model	GPT-4o Judge Metrics $\uparrow$			TOFU Metrics		
	Acc _{forget}	Acc _{recover}	Acc _{retain}	Ext. Strength $\downarrow$	Model Utility $\uparrow$	ROUGE-L _f $\downarrow$
Target	0.19	0.14	0.94	0.99	0.37	0.71
Ideal	0.99	0.99	1.00	0.07	0.38	0.37
MSA _{500B}	0.78 (84.5%)	0.31 69.1%	0.64 68.4%	0.05 +100%	0.41 +100%	0.34 +100%
MSA _{2207B}	0.76 82.1%	0.40 87.8%	0.85 91.2%	0.12 55.8%	0.36 94.2%	0.35 +100%
MSA _{3691B}	0.83 89.9%	0.44 96.7%	0.85 90.6%	0.08 84.1%	0.36 95.9%	0.34 +100%
MSA _{3859B}	0.82 88.9%	0.45 100.0%	0.83 89.0%	0.06 +100%	0.35 93.3%	0.34 +100%
MSA _{last}	0.84 91.6%	0.42 93.9%	0.82 88.0%	0.06 +100%	0.36 93.7%	0.33 +100%
NPO	0.71 77.2%	0.30 66.3%	0.76 81.3%	0.08 84.7%	0.33 86.6%	0.33 +100%
RMU	0.92 100.0%	0.08 17.7%	0.94 100.0%	0.06 +100%	0.37 97.4%	0.14 +100%
GradDiff	0.45 49.2%	0.23 49.7%	0.83 89.0%	0.17 37.3%	0.41 +100%	0.42 87.5%
Task Vector	0.53 57.9%	0.26 57.5%	0.82 87.7%	0.24 27.0%	0.37 97.4%	0.43 87.0%
SatImp	0.28 30.7%	0.17 38.7%	0.90 95.7%	0.40 16.5%	0.37 98.2%	0.55 68.0%
UNDIAL	0.48 52.7%	0.23 50.8%	0.86 92.2%	0.06 +100%	0.39 +100%	0.39 96.0%

ideal

too early

good variants

Every baseline fails on some metric; good MSA variants do not.

MSA on *RESTOR*: Recovery, Not Silence

Llama-3.1-8B, accuracy (%) on 1,051 fact questions

Target	Ideal	MSA		NPO	GradDiff	Task Vec.	SatImp	RMU
44.3	64.8	59.4	64.0	48.5	26.1	44.5	49.2	41.5
		base	instruct					

MSA on MUSE: The Data-Level Metrics

	Ext. Str.↓	Exact Mem↓	Verb Mem↓	Min-k% ↓	Min-k%++ ↓	Know Mem _r ↑	Priv Leak → 0	
Target	.43	.94	.49	1.00	1.00	.62	-100.0	
Ideal	.02	.54	.17	.45	.39	.67	0.0	ideal
MSA _{500B}	.01 ✓	.41 ✓	.12 ✓	.14 ✓	.09 ✓	.51	56.4	too early
MSA _{2207B}	.01 ✓	.37 ✓	.10 ✓	.04 ✓	.01 ✓	.45	74.1	
MSA _{3691B}	.02 ✓	.51 ✓	.15 ✓	.30 ✓	.21 ✓	.63	27.6 ✓	good variants
MSA _{3859B}	.02 ✓	.51 ✓	.15 ✓	.23 ✓	.16 ✓	.59	23.5 ✓	
MSA _{last}	.02	.55	.16 ✓	.37 ✓	.22 ✓	.65 ✓	14.7 ✓	
NPO	.02	.64	.15 ✓	<u>1.00</u>	<u>.99</u>	.62	<u>-99.9</u>	
RMU	.01 ✓	.06 ✓	.08 ✓	.55	.47	.64	-17.8 ✓	
GradDiff	.01 ✓	.20 ✓	.01 ✓	.50	.45	<u>.45</u>	-9.5 ✓	
Task Vector	.01 ✓	.46 ✓	.13 ✓	<u>.92</u>	<u>.95</u>	.48	-84.3	
SatImp	<u>.37</u>	<u>.93</u>	<u>.43</u>	<u>1.00</u>	<u>1.00</u>	.62	-100.0	
UNDIAL	.02	.64	.16 ✓	<u>1.00</u>	<u>1.00</u>	.53	-100.0	

How Old Can the Checkpoint Be?

A gap of **hundreds of billions of tokens** is survivable.

Trillions are not.

What the Two Papers Say Together

- ▶ Concept-level benchmarks reward **silence**. *RESTOR* asks for **recovery**.
- ▶ **Training history is a resource**: a checkpoint from before the data gives a cleaner estimate of what that data did.

Additional Work

Mechanism design for LLM advertising

The Click Is Disappearing



best laptop for a computer science student

AI overview

For coursework and light development, look for at least 16 GB of memory and a recent 8-core processor. Battery life matters more than raw speed. Most students are well served by. . .

[laptopreview.example](#) — 12 best laptops for students

[techstore.example](#) — student laptop deals

The answer arrives **above** the links.

Fewer clicks.

The Click Is Disappearing



best laptop for a computer science student

AI overview

For coursework and light development, look for at least 16 GB of memory and a recent 8-core processor. Battery life matters more than raw speed. Most students are well served by. . .

[laptopreview.example](#) — 12 best laptops for students

[techstore.example](#) — student laptop deals

and search ads were paid **per click**

The answer arrives **above** the links.

Fewer clicks.

The Click Is Disappearing



best laptop for a computer science student

AI overview

For coursework and light development, look for at least 16 GB of memory and a recent 8-core processor. Battery life matters more than raw speed. Most students are well served by. . .

[laptopreview.example](#) — 12 best laptops for students

[techstore.example](#) — student laptop deals

and search ads were paid **per click**

The answer arrives **above** the links.

Fewer clicks.

So what pays for the answer?

Subscriptions Are Not the Whole Answer

ChatGPT was free. Claude Code was free. Serving them is not.

Subscriptions Are Not the Whole Answer

ChatGPT was free. Claude Code was free. Serving them is not.

And there are **free models** now, good enough for most of what people ask.

If a competitor gives it away, a subscription is hard to defend.

Subscriptions Are Not the Whole Answer

ChatGPT was free. Claude Code was free. Serving them is not.

And there are **free models** now, good enough for most of what people ask.

If a competitor gives it away, a subscription is hard to defend.

Search paid for itself with **advertising**.
Can an ad live inside the model's output?

Subscriptions Are Not the Whole Answer

ChatGPT was free. Claude Code was free. Serving them is not.

And there are **free models** now, good enough for most of what people ask.

If a competitor gives it away, a subscription is hard to defend.

Search paid for itself with **advertising**.
Can an ad live inside the model's output?

We asked this in **November 2023**, half a year before AI overviews shipped.

A Question That Could Carry an Ad

Original LLM Output

Prompt: How do CPU and RAM interact within a computer?

Output: The CPU (Central Processing Unit) processes data and instructions, and RAM (Random Access Memory) provides fast, temporary storage for the data the CPU is actively working on, enabling quick access and manipulation of data, resulting in faster and more efficient computing.

Modified Output

Prompt: Can you promote laptop L of brand B in the previous response?

Output: The CPU (Central Processing Unit) processes data and instructions, and RAM (Random Access Memory) provides fast, temporary storage for the data the CPU is actively working on, enabling quick access and manipulation of data, resulting in faster and more efficient computing. Notably, "Laptop L" from "Brand B" is optimized for this synergy, offering a powerful CPU and ample RAM, ensuring a seamless and high-performance computing experience.

A Question That Could Carry an Ad

Original LLM Output

Prompt: How do CPU and RAM interact within a computer?

Output: The CPU (Central Processing Unit) processes data and instructions, and RAM (Random Access Memory) provides fast, temporary storage for the data the CPU is actively working on, enabling quick access and manipulation of data, resulting in faster and more efficient computing.

Modified Output

Prompt: Can you promote laptop L of brand B in the previous response?

Output: The CPU (Central Processing Unit) processes data and instructions, and RAM (Random Access Memory) provides fast, temporary storage for the data the CPU is actively working on, enabling quick access and manipulation of data, resulting in faster and more efficient computing. Notably, "Laptop L" from "Brand B" is optimized for this synergy, offering a powerful CPU and ample RAM, ensuring a seamless and high-performance computing experience.

Three parties have to agree at once:

The advertiser
wants the placement

The user
still wants an answer

The click
has to be worth making

And Not Only in Text

Q: Car recommendation, Ad: **BMW 3 Series**



not personalized




user: 60-year-old woman



user: 60-year-old man

Q: Conference in Florida, Ad: **Travel with SunShine**



not personalized



user: 20-year-old
unmarried living in DC



user: 40-year-old married
living in California

The same question applies to a **generated image**: the car in the answer can be a particular car, the holiday a particular operator.

And Not Only in Text

Q: Car recommendation, Ad: **BMW 3 Series**



not personalized




user: 60-year-old woman



user: 60-year-old man

Q: Conference in Florida, Ad: **Travel with SunShine**



not personalized



user: 20-year-old
unmarried living in DC




user: 40-year-old married
living in California

The same question applies to a **generated image**: the car in the answer can be a particular car, the holiday a particular operator.

And it can be **personalized**, which is where the privacy question starts.

Where the Ad Enters



Modification writes the ad into the text. **Bidding** turns it into a bid. **Prediction** estimates relevance and click-through. **Auction** picks the winner and the price.

Feizi, Hajiaghayi, Rezaei, Shin ^{$\alpha\beta$} . *Online Advertisements with LLMs: Opportunities and Challenges*. **ACM SIGecom Exchanges** 2025.

(^{$\alpha\beta$} alphabetical order)

The Hard Parts

How much

Too many ads and the answer stops being an answer.

Where in the output should one go?

Who

Highest bid, or best fit to the question?

Relevance has to enter the allocation.

The Hard Parts

How much

Too many ads and the answer stops being an answer.

Where in the output should one go?

What it costs

Charge per impression, per click, or per sentence?

Truthful bidding must stay the best strategy.

Who


Highest bid, or best fit to the question?

Relevance has to enter the allocation.

And the rest

Privacy of the user context. Latency of the auction. Advertisers who misbehave.

Ad Auctions via Retrieval



One auction per **segment**. An ad wins with probability set by its **bid** and its **relevance**, then joins the context and the model writes the next segment around it.

Hajiaghayi, Lahaie, Rezaei, Shin ^{$\alpha\beta$} . *Ad Auctions for LLMs via Retrieval Augmented Generation*. **NeurIPS 2024**. (^{$\alpha\beta$} alphabetical order)

How the Winner Is Picked

Segment auction

For each segment:

$$\text{score } s_i = \underbrace{q_i}_{\text{relevance}} \cdot \underbrace{b_i}_{\text{bid}} \cdot e^{\varepsilon_i}, \quad \varepsilon_i \sim \text{Gumbel}$$

the winner is $\arg \max_i s_i$

it pays the **lowest bid that would still have won**, charged **per click**

How the Winner Is Picked

Segment auction

For each segment:

$$\text{score } s_i = \underbrace{q_i}_{\text{relevance}} \cdot \underbrace{b_i}_{\text{bid}} \cdot e^{\varepsilon_i}, \quad \varepsilon_i \sim \text{Gumbel}$$

the winner is $\arg \max_i s_i$

it pays the **lowest bid that would still have won**, charged **per click**

The noise makes it a **lottery weighted by $q_i b_i$** : a highly relevant ad can beat a higher bid.

How the Winner Is Picked

Segment auction

For each segment:

$$\text{score } s_i = \underbrace{q_i}_{\text{relevance}} \cdot \underbrace{b_i}_{\text{bid}} \cdot e^{\varepsilon_i}, \quad \varepsilon_i \sim \text{Gumbel}$$

the winner is $\arg \max_i s_i$

it pays the **lowest bid that would still have won**, charged **per click**

The noise makes it a **lottery weighted by $q_i b_i$** : a highly relevant ad can beat a higher bid.

The price is the second-price idea, so **bidding your true value is a dominant strategy**.

What the Mechanism Guarantees

- ▶ **Truthful** (DSIC). No advertiser gains by misreporting its value.
- ▶ **Individually rational**. Nobody is worse off for taking part.

What the Mechanism Guarantees

- ▶ **Truthful** (DSIC). No advertiser gains by misreporting its value.
- ▶ **Individually rational**. Nobody is worse off for taking part.
- ▶ **Maximal logarithmic social welfare** $\prod_i x_i^{v_i q_i}$ among independent segment auctions.

What the Mechanism Guarantees

- ▶ **Truthful** (DSIC). No advertiser gains by misreporting its value.
- ▶ **Individually rational**. Nobody is worse off for taking part.
- ▶ **Maximal logarithmic social welfare** $\prod_i x_i^{v_i q_i}$ among independent segment auctions.
- ▶ A **log** welfare, not a sum. It rewards spreading allocation across advertisers rather than handing every segment to one, so efficiency and fairness come together.

What the Mechanism Guarantees

- ▶ **Truthful** (DSIC). No advertiser gains by misreporting its value.
- ▶ **Individually rational**. Nobody is worse off for taking part.
- ▶ **Maximal logarithmic social welfare** $\prod_i x_i^{v_i q_i}$ among independent segment auctions.
- ▶ A **log** welfare, not a sum. It rewards spreading allocation across advertisers rather than handing every segment to one, so efficiency and fairness come together.

Ignoring relevance earns more revenue and serves the user worse.

This thesis asks where a foundation model's knowledge comes from, and what we can do once we know.

- ▶ **Where knowledge lives.** Methods and heuristics for localizing a concept inside a generative model, and for editing it in place.

This thesis asks where a foundation model's knowledge comes from, and what we can do once we know.

- ▶ **Where knowledge lives.** Methods and heuristics for localizing a concept inside a generative model, and for editing it in place.
- ▶ **What the data did.** Tools for studying the effect of training data through machine unlearning: a benchmark that asks for recovery, and a method that reads the training history.

This thesis asks where a foundation model's knowledge comes from, and what we can do once we know.

- ▶ **Where knowledge lives.** Methods and heuristics for localizing a concept inside a generative model, and for editing it in place.
- ▶ **What the data did.** Tools for studying the effect of training data through machine unlearning: a benchmark that asks for recovery, and a method that reads the training history.
- ▶ **How it sustains itself.** Mechanism design for generative AI: auctions that place advertising inside a model's output, with the incentive guarantees such a market needs.

Conclusion

This thesis asks where a foundation model's knowledge comes from, and what we can do once we know.

- ▶ **Where knowledge lives.** Methods and heuristics for localizing a concept inside a generative model, and for editing it in place.
- ▶ **What the data did.** Tools for studying the effect of training data through machine unlearning: a benchmark that asks for recovery, and a method that reads the training history.
- ▶ **How it sustains itself.** Mechanism design for generative AI: auctions that place advertising inside a model's output, with the incentive guarantees such a market needs.

**Treat the model and its data as objects we can inspect and change,
not as a finished artifact.**

Other Work: Interpretability and Unlearning

- ▶ Revisiting the Past: Data Unlearning with Model State History. **ICLR 2026.**
K. Rezaei, M. Saberi, A. Ravichander*, S. Feizi*
- ▶ *RESTOR*: Knowledge Recovery through Machine Unlearning. **TMLR 2025.**
K. Rezaei, K. Chandu, S. Feizi, Y. Choi, F. Brahman, A. Ravichander
- ▶ Localizing Knowledge in Diffusion Transformers. **NeurIPS 2025.**
A. Zarei, S. Basu, K. Rezaei, Z. Lin, S. Nag, S. Feizi
- ▶ On Mechanistic Knowledge Localization in Text-to-Image Generative Models. **ICML 2024.**
S. Basu*, K. Rezaei*, R. Rossi, C. Zhao, V. Morariu, V. Manjunatha, S. Feizi
- ▶ PRIME: Prioritizing Interpretability in Failure Mode Extraction. **ICLR 2024.**
K. Rezaei*, M. Saberi*, M. Moayeri, S. Feizi
- ▶ Text-To-Concept (and Back) via Cross-Model Alignment. **ICML 2023.**
M. Moayeri*, K. Rezaei*, M. Sanjabi, S. Feizi
- ▶ Improving Compositional Attribute Binding in Text-to-Image Models. **Preprint.**
A. Zarei*, K. Rezaei*, S. Basu, M. Saberi, M. Moayeri, P. Kattakinda, S. Feizi

Other Work: LLM Reasoning and Robustness

LLM reasoning

- ▶ Reasoning Under Uncertainty: Exploring Probabilistic Reasoning Capabilities of LLMs. **EMNLP 2026**.
M. Pournemat, **K. Rezaei**, G. Sriramanan, A. Zarei, J. Fu, Y. Wang, H. Eghbalzadeh, S. Feizi
- ▶ RePanda: Pandas-powered Tabular Verification and Reasoning. **ACL 2025**.
A. Chegini, **K. Rezaei**, H. Eghbalzadeh, S. Feizi

Robustness

- ▶ Run-Off Election: Improved Provable Defense against Data Poisoning Attacks. **ICML 2023**.
K. Rezaei*, K. Banihashem*, A. Chegini, S. Feizi
- ▶ Robustness of AI-Image Detectors: Fundamental Limits and Practical Attacks. **ICLR 2024**.
M. Saberi, V. S. Sadasivan, **K. Rezaei**, A. Kumar, A. Chegini, W. Wang, S. Feizi

Other Work: Mechanism Design

- ▶ Ad Auctions for LLMs via Retrieval Augmented Generation. **NeurIPS 2024**.
M. Hajiaghayi, S. Lahaie, **K. Rezaei**, S. Shin
- ▶ Online Advertisements with LLMs: Opportunities and Challenges. **SIGecom Exchanges 2025**.
S. Feizi, M. Hajiaghayi, **K. Rezaei**, S. Shin
- ▶ Multi-agent Delegated Search. **EC 2023**.
M. Hajiaghayi, **K. Rezaei**, S. Shin
- ▶ Regret Analysis of Repeated Delegated Choice. **AAAI 2024**.
M. Hajiaghayi, M. Mahdavi, **K. Rezaei**, S. Shin

Where the Work Was Done



Reliable AI Lab, University of Maryland

Fall 2022 – Summer 2026



Susquehanna International Group

Summer 2026



Adobe Research

Summer 2025



Google Research

Spring 2025



Allen Institute for AI

Summer 2024

With Thanks

Soheil Feizi, for the opportunity to work on hard and exciting problems, and for years of guidance.

MohammadTaghi Hajiaghayi, for a second research world to explore and for the theory behind it.

Maria K. Cameron, **David Jacobs** and **Sarah Wiegrefe**, for serving on the committee and for their feedback on this dissertation.

Abhilasha Ravichander, for her mentorship. **Mehrdad Saberi**, **Mazda Moayeri**, **Samyadeep Basu**, **Arman Zarei** and **Suho Shin**, for collaborations that made the work better and the years more enjoyable.

Colleagues and mentors at **Google Research**, **Adobe Research**, **Susquehanna International Group** and the **Allen Institute for AI**.

And my **father**, **mother** and **sister**, who carried me through, and my **friends**, who kept the years grounded.

With Thanks

Soheil Feizi, for the opportunity to work on hard and exciting problems, and for years of guidance.

MohammadTaghi Hajiaghayi, for a second research world to explore and for the theory behind it.

Maria K. Cameron, **David Jacobs** and **Sarah Wiegrefe**, for serving on the committee and for their feedback on this dissertation.

Abhilasha Ravichander, for her mentorship. **Mehrdad Saberi**, **Mazda Moayeri**, **Samyadeep Basu**, **Arman Zarei** and **Suho Shin**, for collaborations that made the work better and the years more enjoyable.

Colleagues and mentors at **Google Research**, **Adobe Research**, **Susquehanna International Group** and the **Allen Institute for AI**.

And my **father**, **mother** and **sister**, who carried me through, and my **friends**, who kept the years grounded.


*P.S. Thanks to **Claude Code** for the slides, and to **Gemini** and **ChatGPT** for help with writing the papers. :-)*

Thank you

for your attention

Questions?


Some Concepts Are Tighter Than Others



Intervening on the top- $K\%$ of blocks. Copyright, place and celebrity collapse early; animals are spread much wider.





DreamBooth: One Subject, One New Token

DreamBooth: a few photos teach the token **[V]**.




Localized tuning still obeys “*origami*”; full tuning drifts back to the photos.

Personalization: Better on Every Metric



- ▶ Higher prompt alignment on the customized subject.
- ▶ Surrounding classes (Husky, Golden Retriever) keep their identity.
- ▶ FID holds flat while full fine-tuning degrades with more steps.
- ▶ All of it while touching **9 of 28 blocks**.

Unlearning: Same Erasure, Less Collateral



- ▶ Target style erased as fast as full fine-tuning.
- ▶ Surrounding artists keep their identity.
- ▶ Alignment with the anchor is preserved.
- ▶ FID degrades far less; only **5 of 28 blocks** move.

References I

- [Bas+24] Samyadeep Basu et al. “On Mechanistic Knowledge Localization in Text-to-Image Generative Models”. In: *International Conference on Machine Learning (ICML)*. 2024.
- [Che+23] Junsong Chen et al. “Pixart- α : Fast training of diffusion transformer for photorealistic text-to-image synthesis”. In: *arXiv preprint arXiv:2310.00426* (2023).
- [Fei+25] Soheil Feizi et al. “Online Advertisements with LLMs: Opportunities and Challenges”. In: *ACM SIGecom Exchanges* (2025).
- [GAS20] Aditya Golatkar, Alessandro Achille, and Stefano Soatto. “Eternal sunshine of the spotless net: Selective forgetting in deep networks”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020, pp. 9304–9312.
- [Gev+21] Mor Geva et al. “Transformer Feed-Forward Layers Are Key-Value Memories”. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 2021.
- [Haj+24] MohammadTaghi Hajiaghayi et al. “Ad Auctions for LLMs via Retrieval Augmented Generation”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2024.
- [Jan+22] Joel Jang et al. “Knowledge unlearning for mitigating privacy risks in language models”. In: *arXiv preprint arXiv:2210.01504* (2022).
- [Kum+23] Nupur Kumari et al. “Ablating concepts in text-to-image diffusion models”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision*. 2023, pp. 22691–22702.

References II

- [LPG+24] Nathaniel Li, Alexander Pan, Anjali Gopal, et al. “The WMDP Benchmark: Measuring and Reducing Malicious Use with Unlearning”. In: *arXiv preprint arXiv:2403.03218* (2024).
- [Mai+24] Pratyush Maini et al. “Tofu: A task of fictitious unlearning for llms”. In: *arXiv preprint arXiv:2401.06121* (2024).
- [Men+22a] Kevin Meng et al. “Locating and Editing Factual Associations in GPT”. In: *Advances in Neural Information Processing Systems* 36 (2022). arXiv:2202.05262.
- [Men+22b] Kevin Meng et al. “Mass-editing memory in a transformer”. In: *arXiv preprint arXiv:2210.07229* (2022).
- [Rad+21] Alec Radford et al. “Learning transferable visual models from natural language supervision”. In: *International conference on machine learning*. PmLR. 2021, pp. 8748–8763.
- [Rez+25] Keivan Rezaei et al. “RESTOR: Knowledge Recovery through Machine Unlearning”. In: *Transactions on Machine Learning Research (TMLR)* (2025).
- [Rez+26] Keivan Rezaei et al. “Revisiting the Past: Data Unlearning with Model State History”. In: *International Conference on Learning Representations (ICLR)*. 2026.
- [Rom+22] Robin Rombach et al. “High-resolution image synthesis with latent diffusion models”. In: *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. 2022, pp. 10684–10695.
- [Rui+23] Nataniel Ruiz et al. “Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation”. In: *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*. 2023, pp. 22500–22510.

References III

- [Shi+24] Weijia Shi et al. “Muse: Machine unlearning six-way evaluation for language models”. In: *arXiv preprint arXiv:2407.06460* (2024).
- [Som+24] Gowthami Somepalli et al. “Measuring style similarity in diffusion models”. In: *arXiv preprint arXiv:2404.01292* (2024).
- [YXL23] Yuanshun Yao, Xiaojun Xu, and Yang Liu. “Large language model unlearning”. In: *arXiv preprint arXiv:2310.10683* (2023).
- [Zar+25] Arman Zarei et al. “Localizing Knowledge in Diffusion Transformers”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. 2025.
- [Zha+24] Ruiqi Zhang et al. “Negative preference optimization: From catastrophic collapse to effective unlearning”. In: *arXiv preprint arXiv:2404.05868* (2024).